

---

IN1011

# Operating Systems

---

## Lecture 02: Processes

# Today's Questions




- What happens when a program's execution is interrupted?
- How does the OS interleave the execution of multiple programs?
- How does the OS ensure new processes are not waiting while the CPU is idle (faster response times)?
- How can limited RAM be shared among interleaved processes?
- What does a process look like to an OS?
- What does an OS do when creating a process?
- Is the OS a process?




# Questions

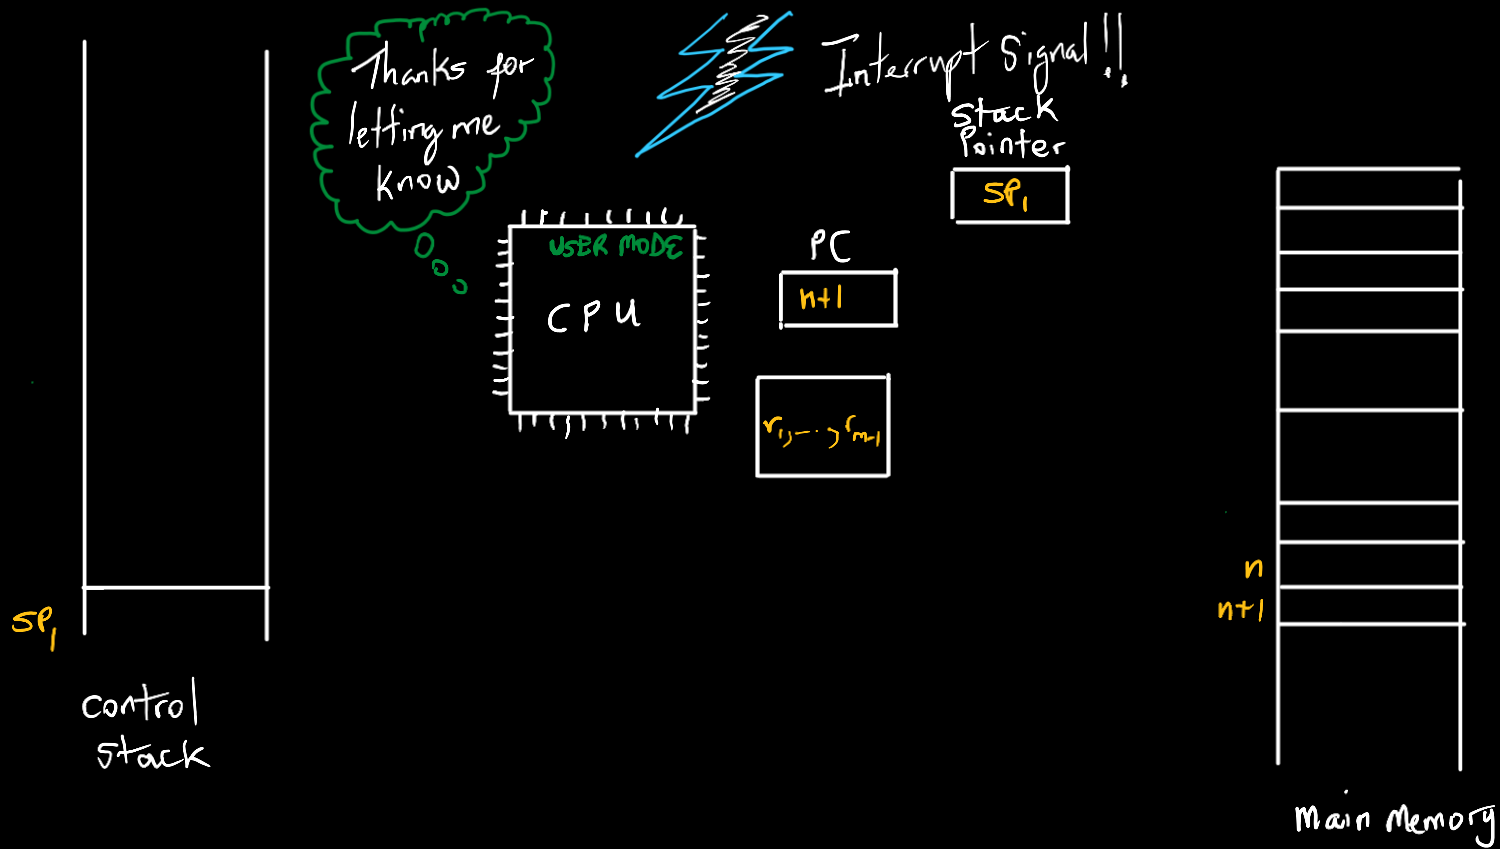
- What happens when a program's execution is interrupted ?




# “Interrupting” Processes

Mechanism	Cause	Use
Interrupt	External to the execution of the current instruction	Reaction to an asynchronous external event
Trap	Associated with the execution of the current instruction	Handling of an error or an exception condition
Supervisor call	Explicit request	Call to an operating system function






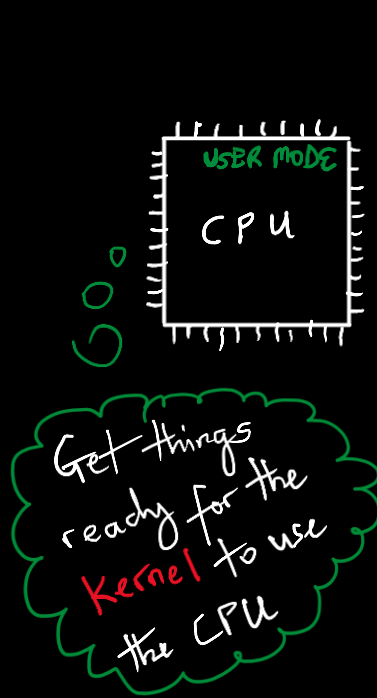









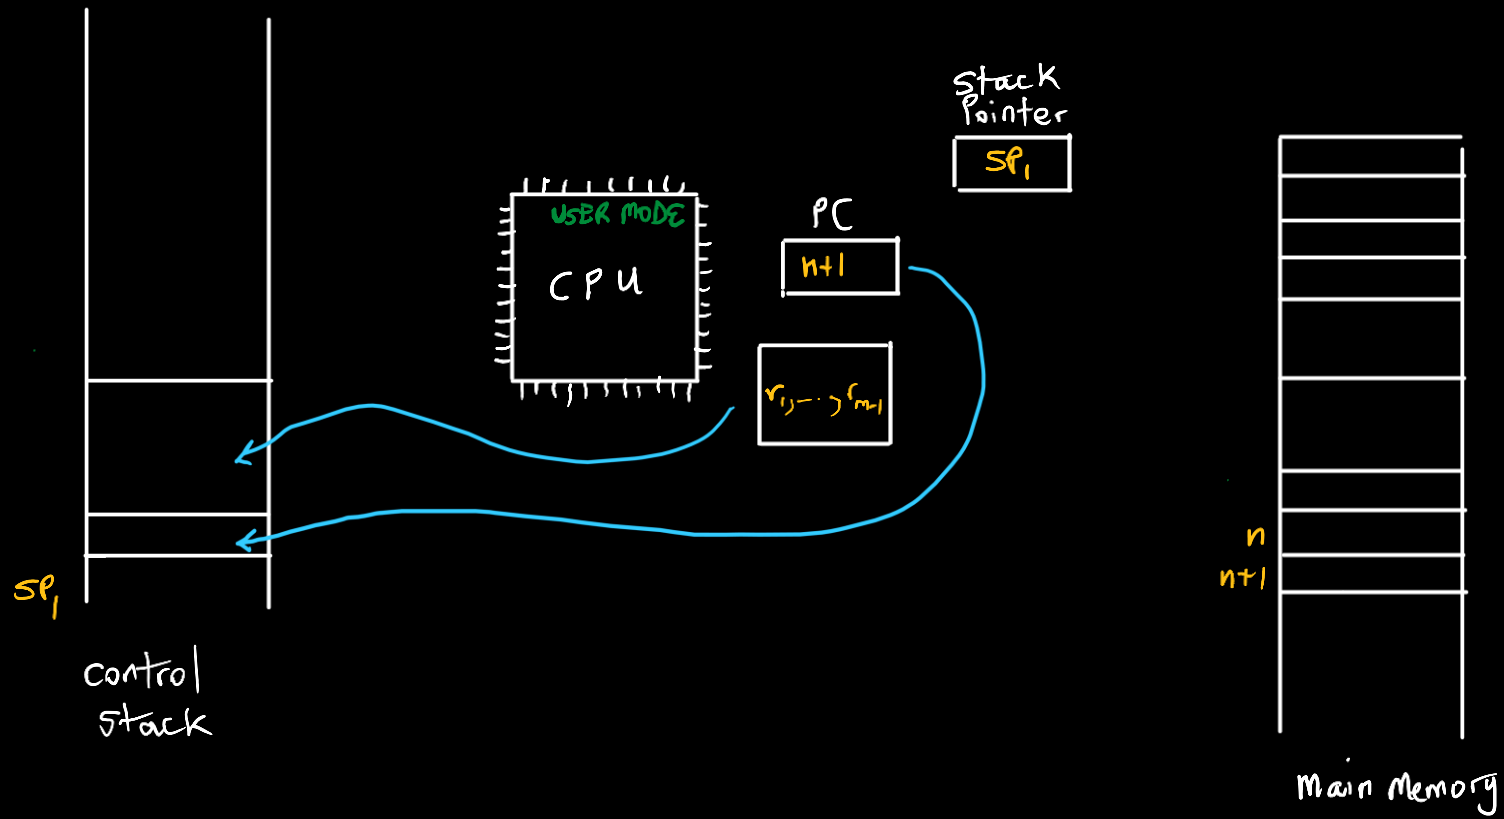


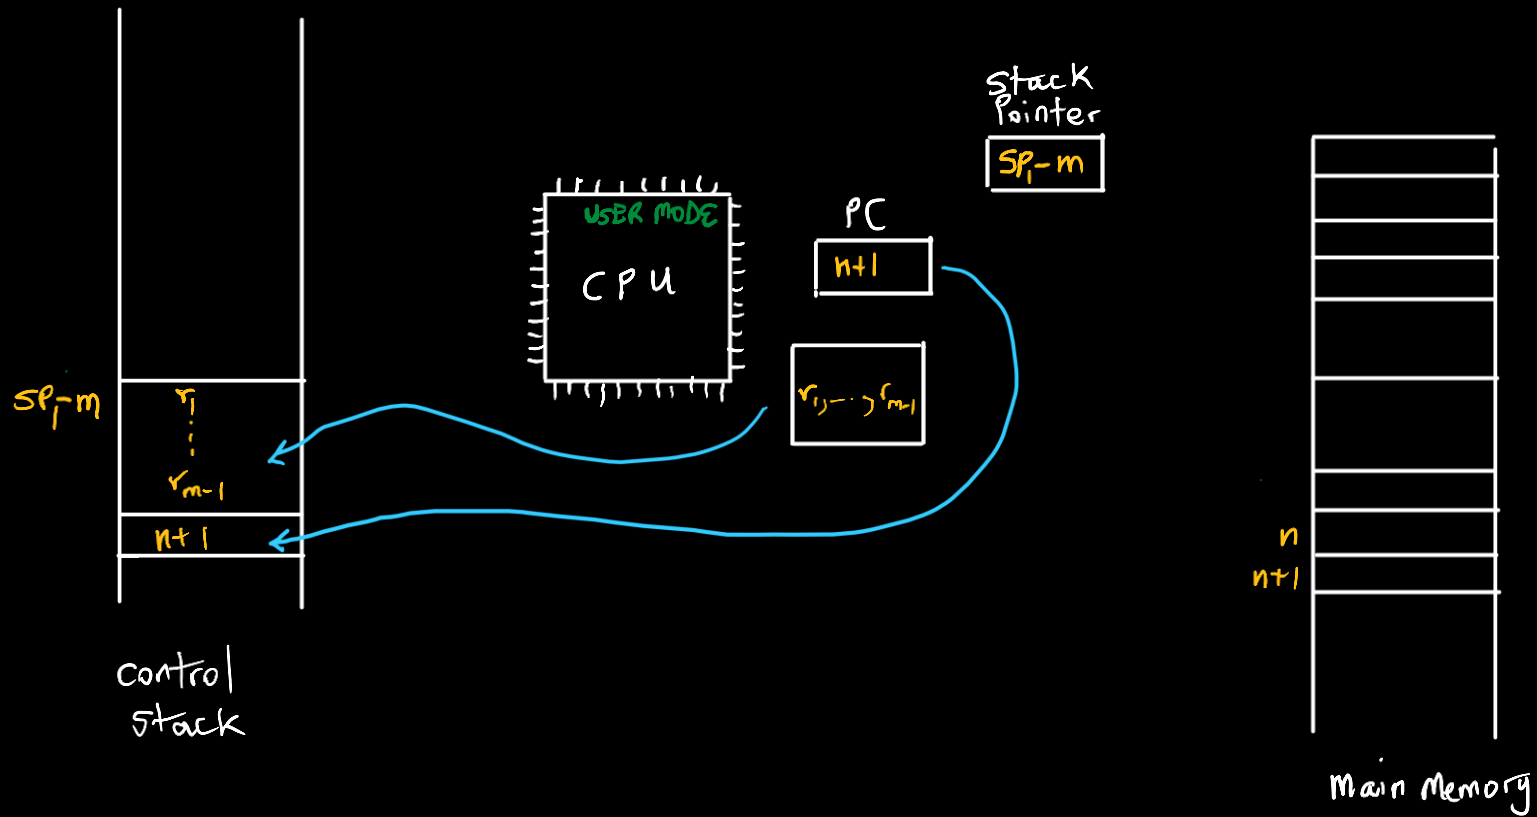























control  
stack




Stack  
Pointer  
 $SP_1-m$







Main Memory







control  
stack








control  
stack



Main Memory




control  
stack




Stack  
Pointer  
 $SP-m$



from kernel's interrupt  
vector




Main Memory



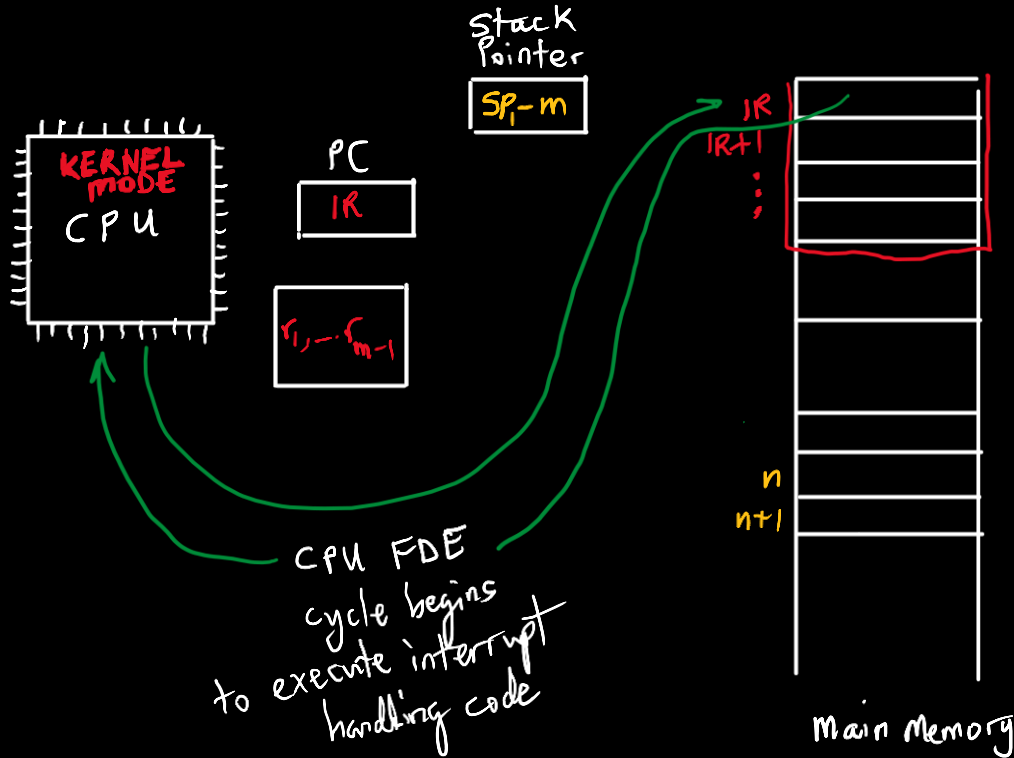
control  
stack



Main Memory




control  
stack




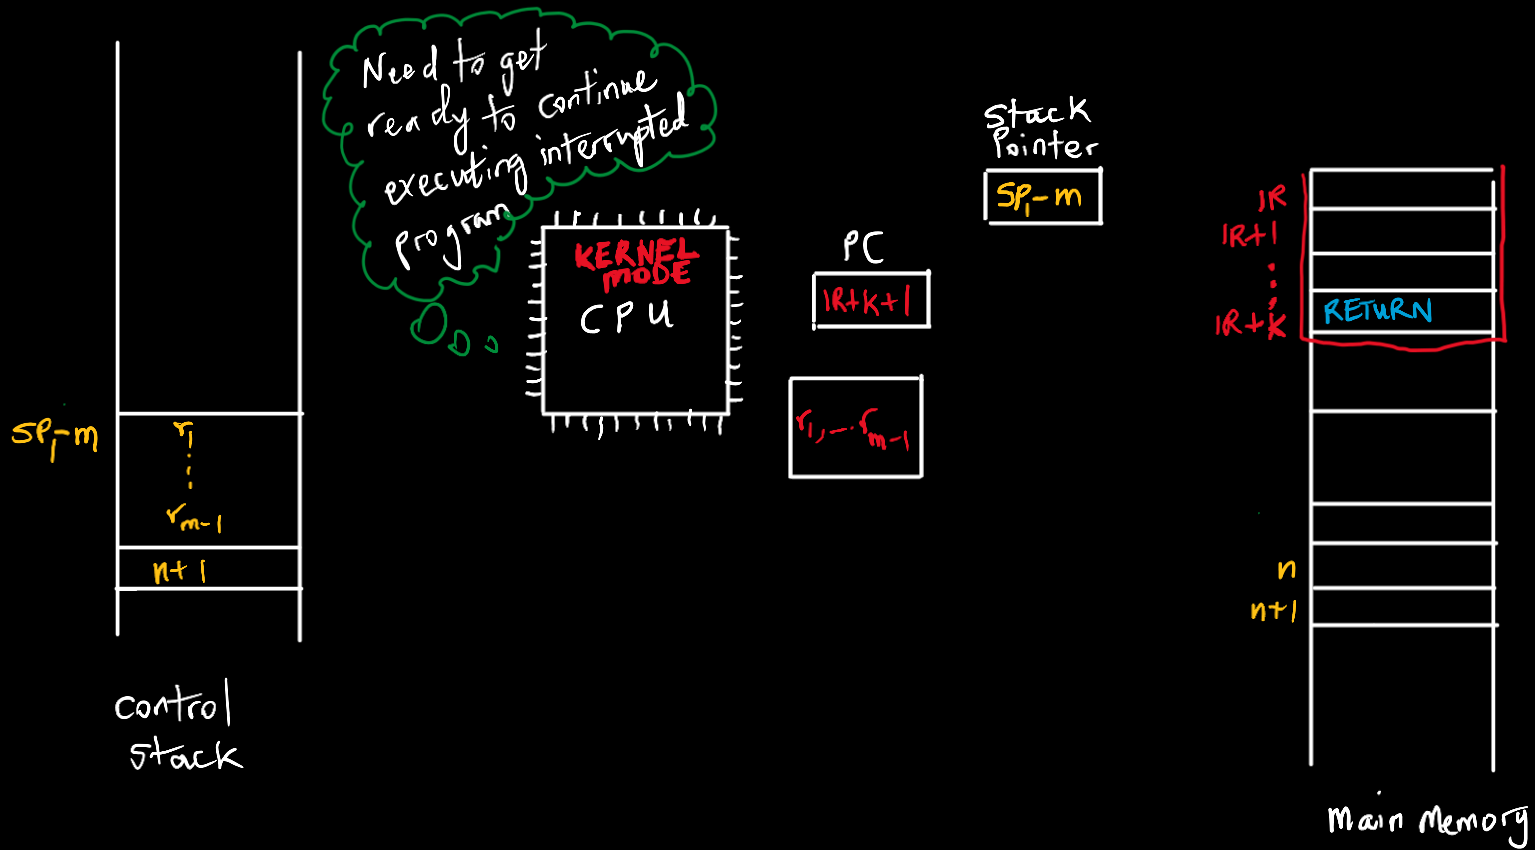
Main Memory

... Sometime Later ...




control  
stack













control  
stack


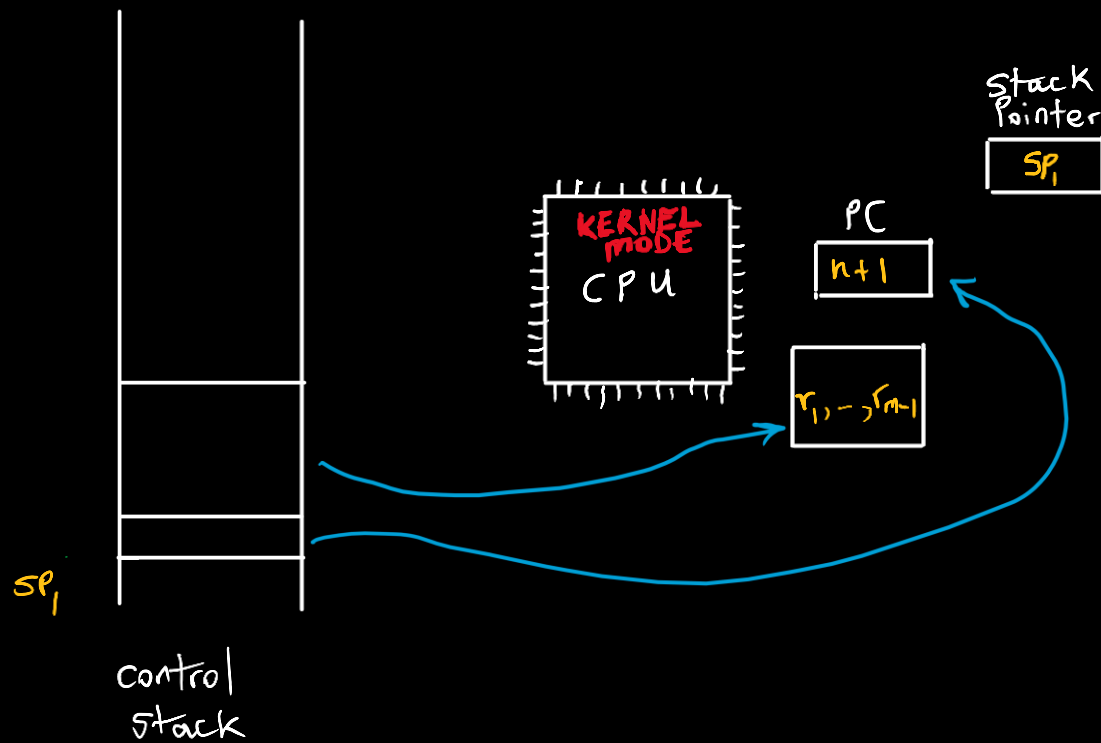





Stack  
Pointer  
 $SP-m$





Main Memory






















# Questions

- How does the OS manage the executions of programs?
  - **partial ans:** OS manages **processes**
- But,...what is a process?

# Definition of a Process

- A process is:
  - a program in execution
  - an instance of a program running on a computer
  - the entity that can be assigned to, and executed on, a processor
  - An executing set of related machine instructions
  - A unit of activity characterised by a single sequential thread of execution, a current state, and an associated set of system resources
- Can be useful to think of programs as being passive, while processes are active

# “Passive” vs “Active”

- A musical score (i.e. sheet music) is “passive”, while the music produced by an orchestra playing the score is “active”
- A cooking recipe is “passive”, while cooking with the recipe is “active”
- A film script is “passive”, while performers acting based on the script is “active”
- In essence, the work done by following “passive” instructions is regarded as “active”

# Management of Program Execution: Processes

- Processes were first used by designers of the Multics OS in the 1960s
- They are the most important abstraction provided by all modern OSes
- The OS uses processes to ensure:
  - Resources are made available to multiple applications
  - The CPU is switched among multiple applications, i.e. each application will appear to progress
  - The CPU and I/O devices can be used efficiently, securely and reliably

# Process Uses Program Elements

Two essential program elements used by a process are:

## Program code

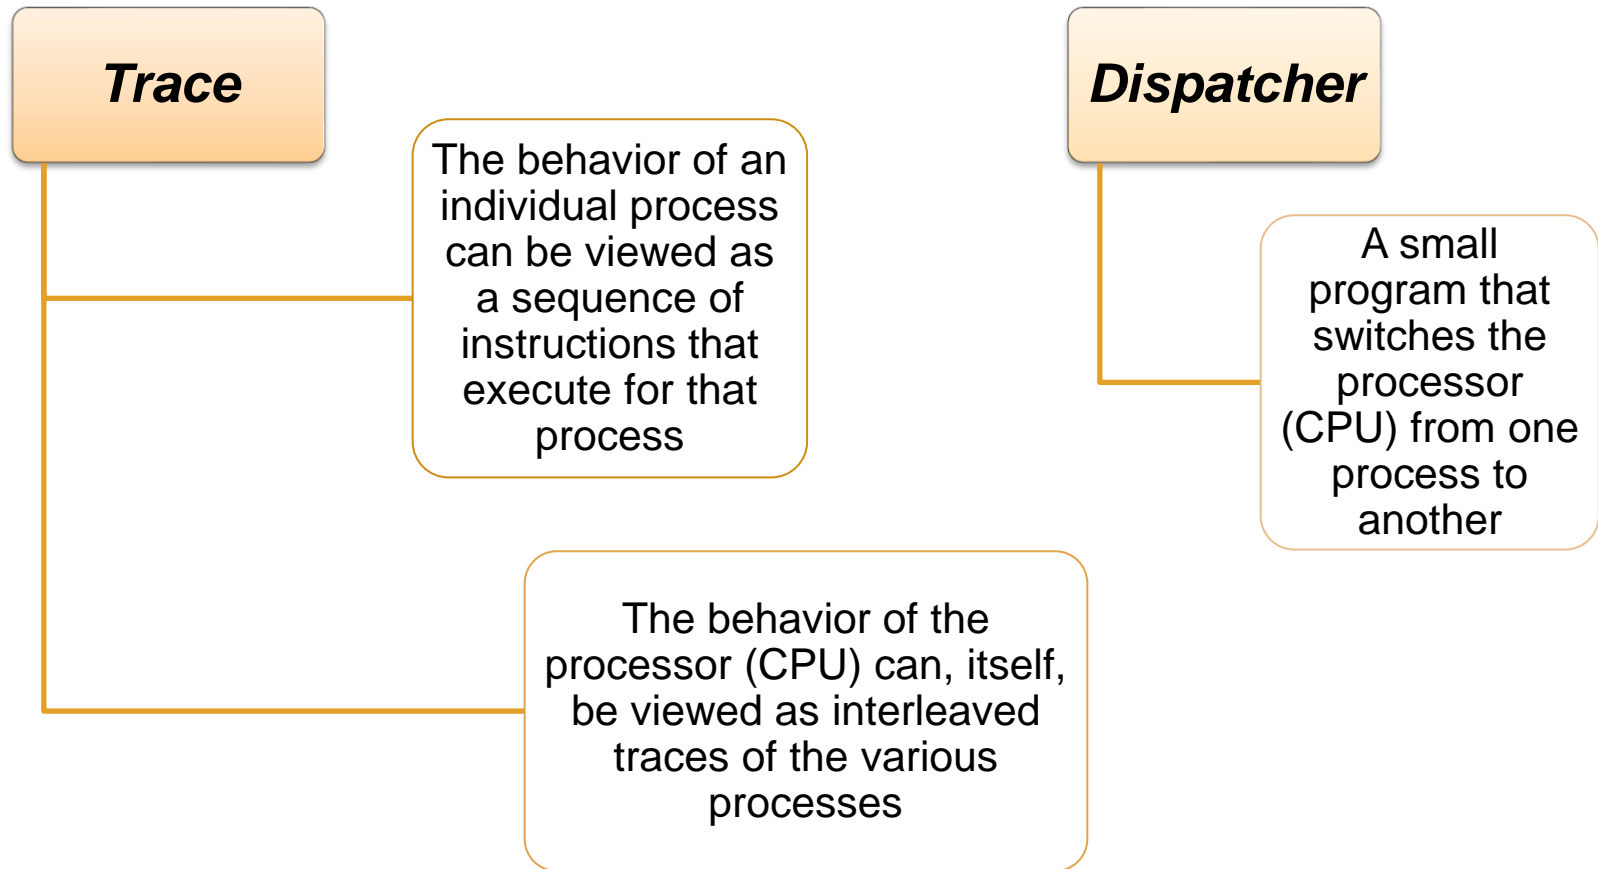
- which may be shared with other processes that are executing the same program

A set of data associated with that code

# Questions

- How does the OS interleave the execution of multiple programs?
- More specifically, how does the OS re-assign hardware resources among multiple programs?

# Interleaving Processes



# Interleaving Processes

More about the **dispatcher** when we cover the topic of *scheduling*

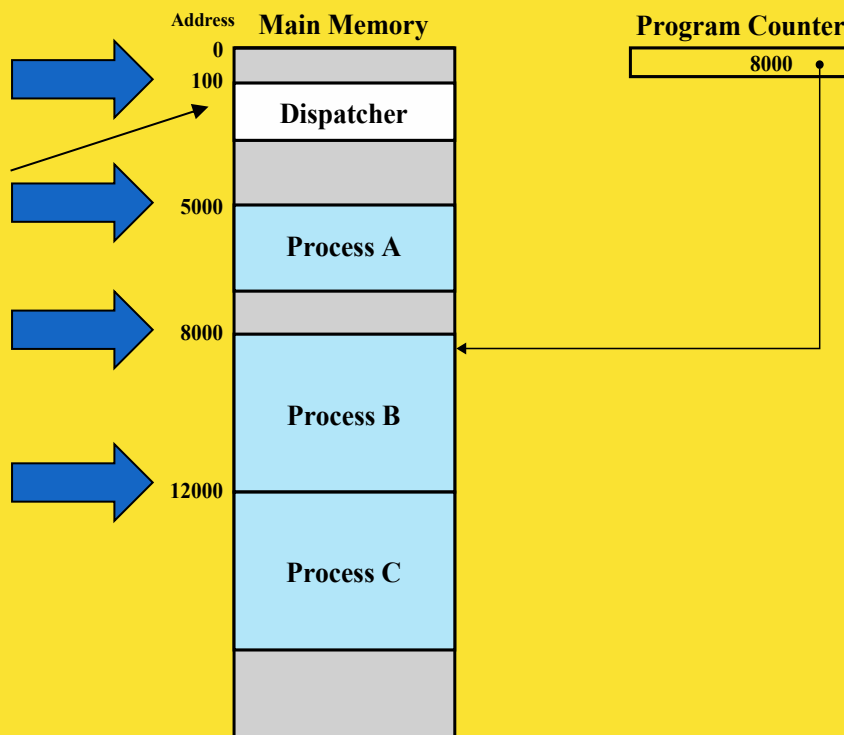


Figure 3.2 Snapshot of Example Execution (Figure 3.4)  
at Instruction Cycle 13



# Interleaving Processes: Each Program's Viewpoint

5000	8000	12000
5001	8001	12001
5002	8002	12002
5003	8003	12003
5004		12004
5005		12005
5006		12006
5007		12007
5008		12008
5009		12009
5010		12010
5011		12011

**(a) Trace of Process A**

**(b) Trace of Process B**

**(c) Trace of Process C**

5000 = Starting address of program of Process A

8000 = Starting address of program of Process B

12000 = Starting address of program of Process C

**Figure 3.3 Traces of Processes of Figure 3.2**

# Interleaving Processes: CPU's Viewpoint

1	5000			27	12004
2	5001			28	12005
3	5002				----- Timeout
4	5003			29	100
5	5004			30	101
6	5005			31	102
				32	103
				33	104
				34	105
				35	5006
				36	5007
				37	5008
				38	5009
				39	5010
				40	5011
					----- Timeout
				41	100
				42	101
				43	102
				44	103
				45	104
				46	105
				47	12006
				48	12007
				49	12008
				50	12009
				51	12010
				52	12011
					----- Timeout

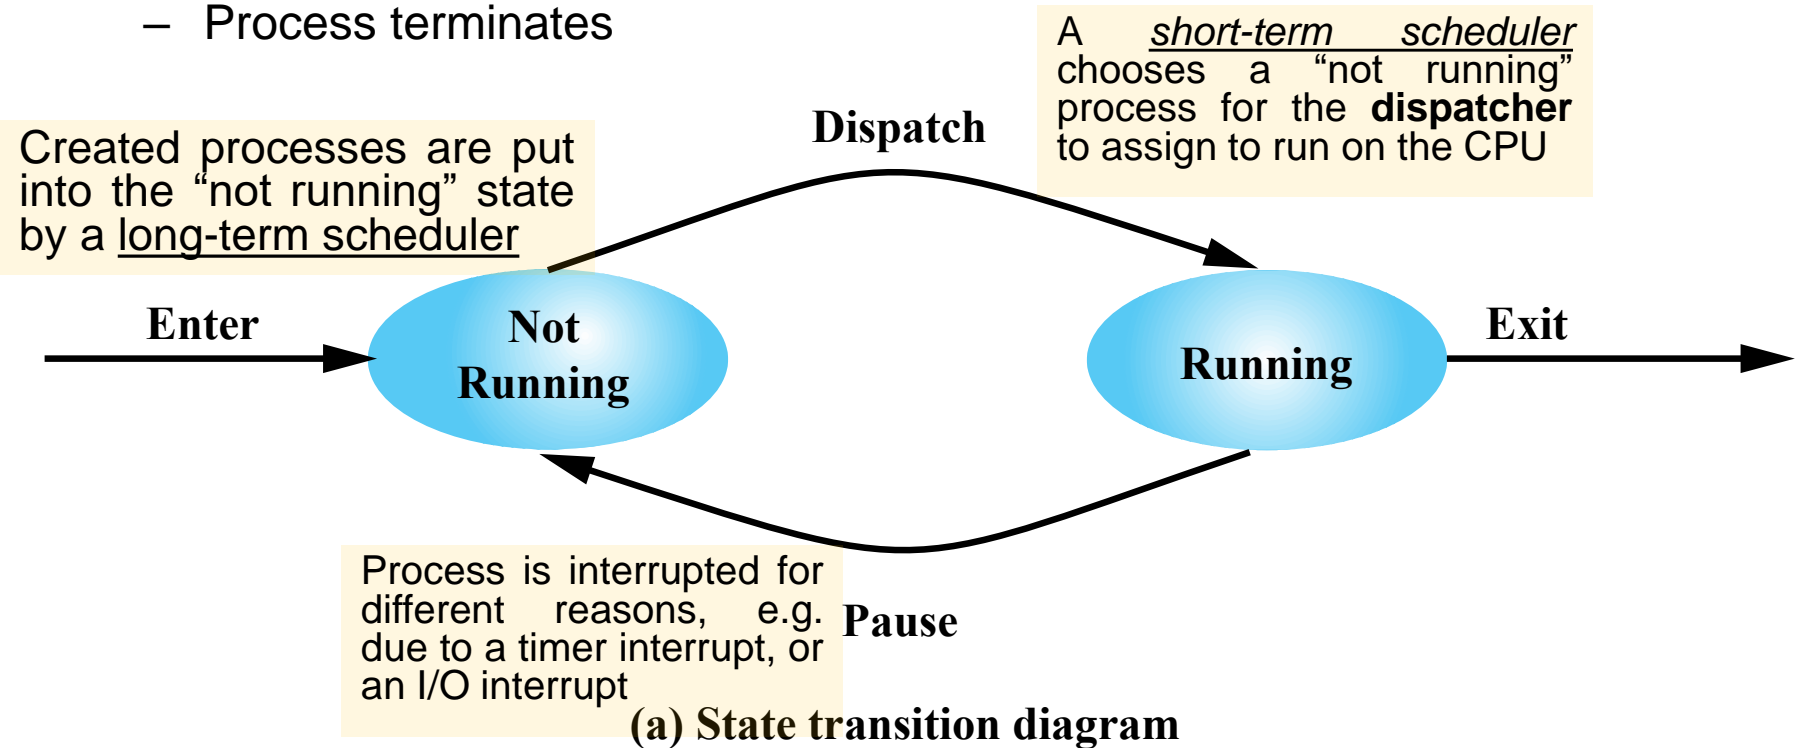
100 = Starting address of dispatcher program

Shaded areas indicate execution of dispatcher process;  
first and third columns count instruction cycles;  
second and fourth columns show address of instruction being executed

Figure 3.4 Combined Trace of Processes of Figure 3.2

# Two-State Process Model

- Simplest model:
  - Process is created
  - Process is always either running or not running
  - Process terminates



# Reasons for Process Creation

- New batch job
- User logon
- OS services
- Spawned by existing process

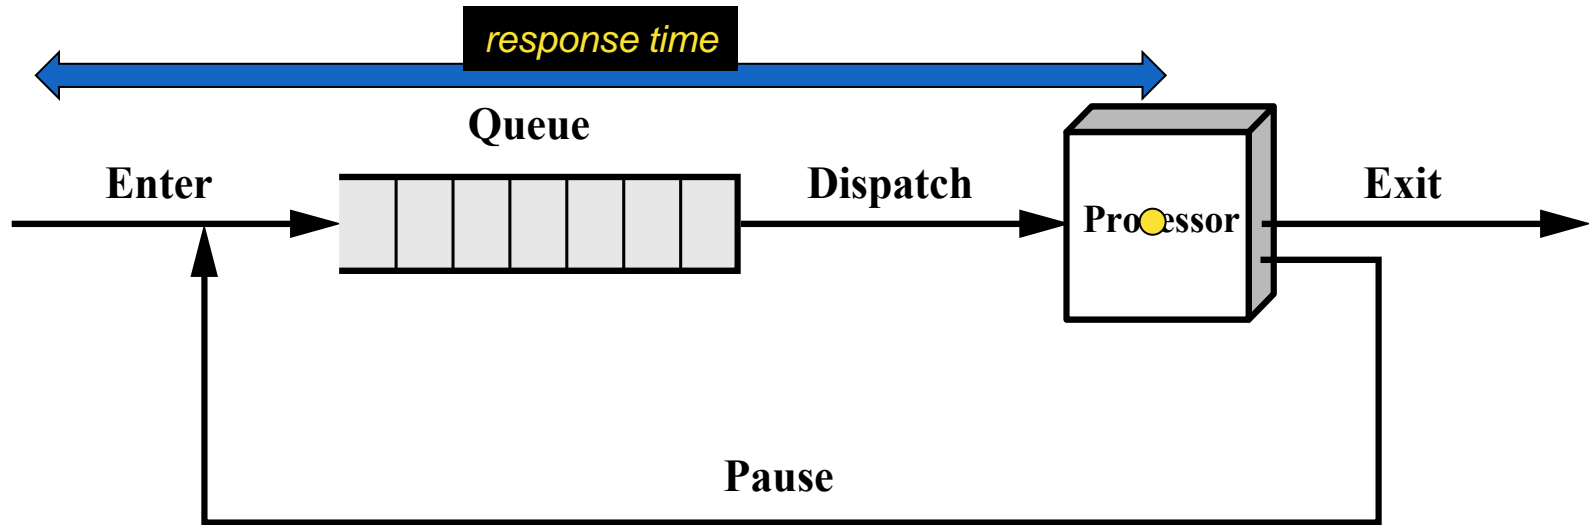
# Reasons for Process Termination

- Normal completion
- Time limit
- Insufficient computer resource available (e.g. RAM)
- Significant errors
  - I/O
  - division by zero
  - Arithmetic
- Parent process request or termination
- User termination
- Privileged instruction

# Question

- How can an OS interleave processes so that new processes are not waiting while the CPU is idle?

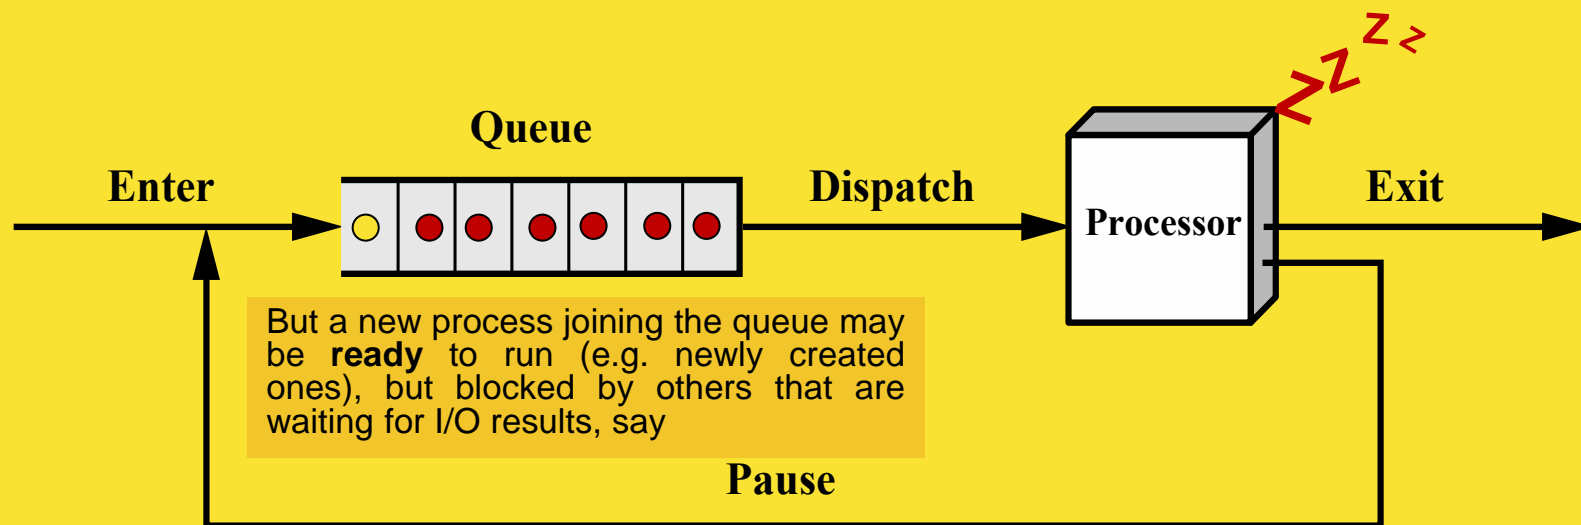
# Response Time



(b) Queuing diagram

Figure 3.5 Two-State Process Model

# But ...



(b) Queuing diagram

Figure 3.5 Two-State Process Model



# Five-State Process Model

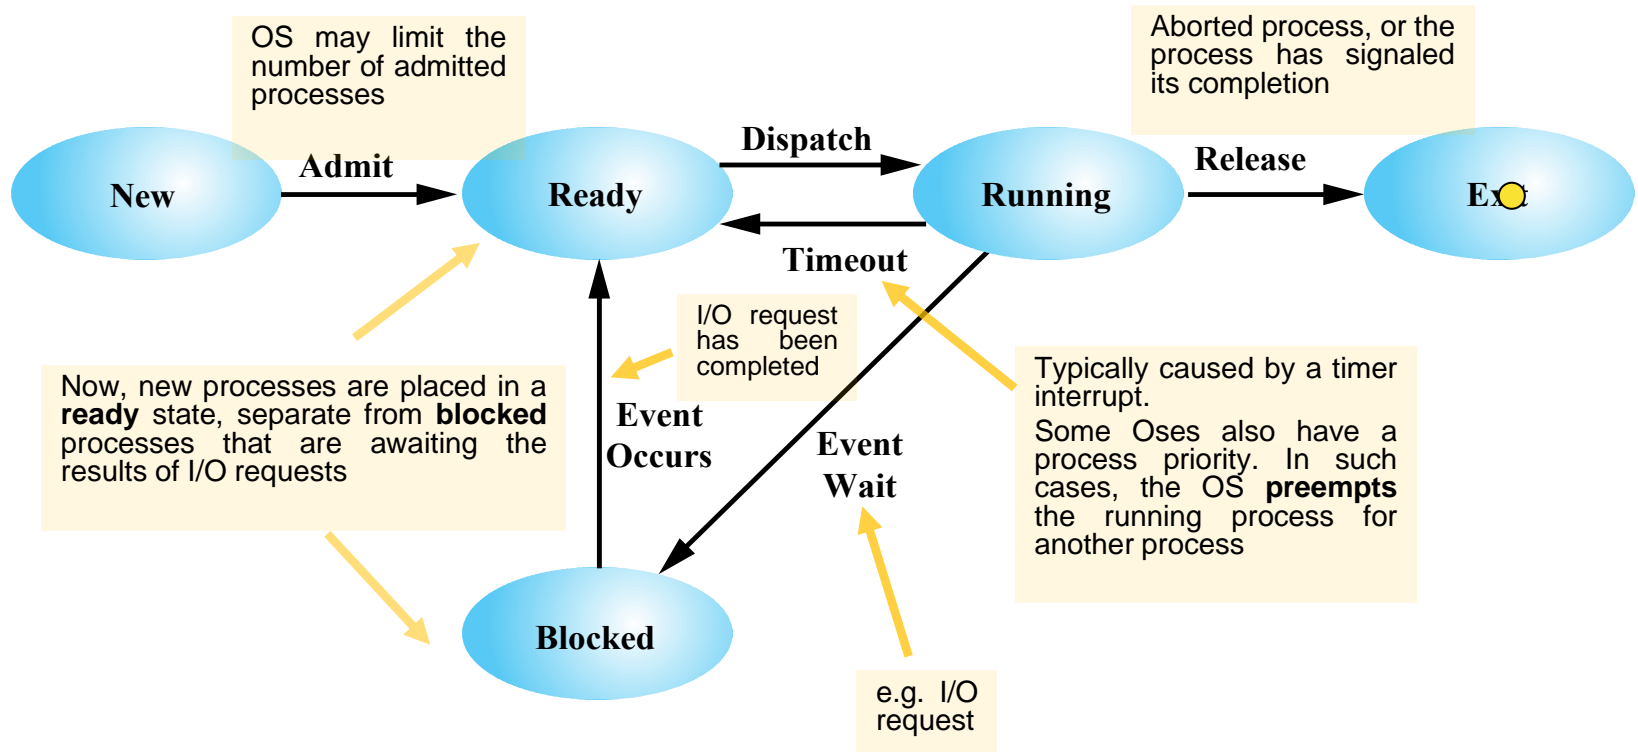

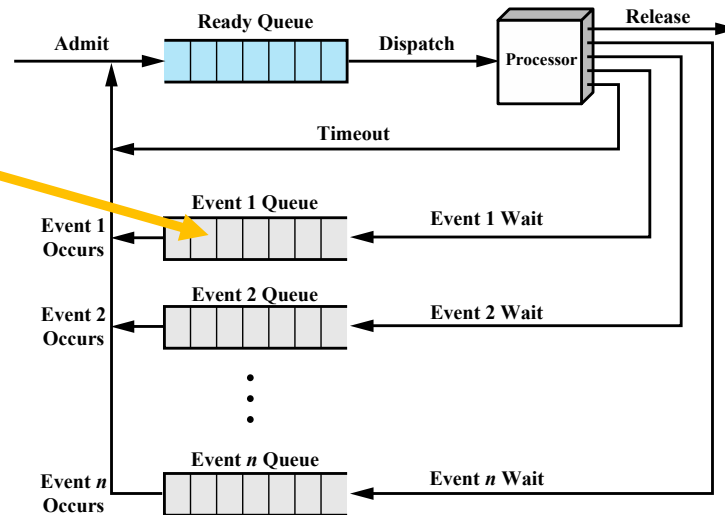


Figure 3.6 Five-State Process Model



(a) Single blocked queue

Within each queue, different processes may be awaiting results from requests made to the same I/O device, say

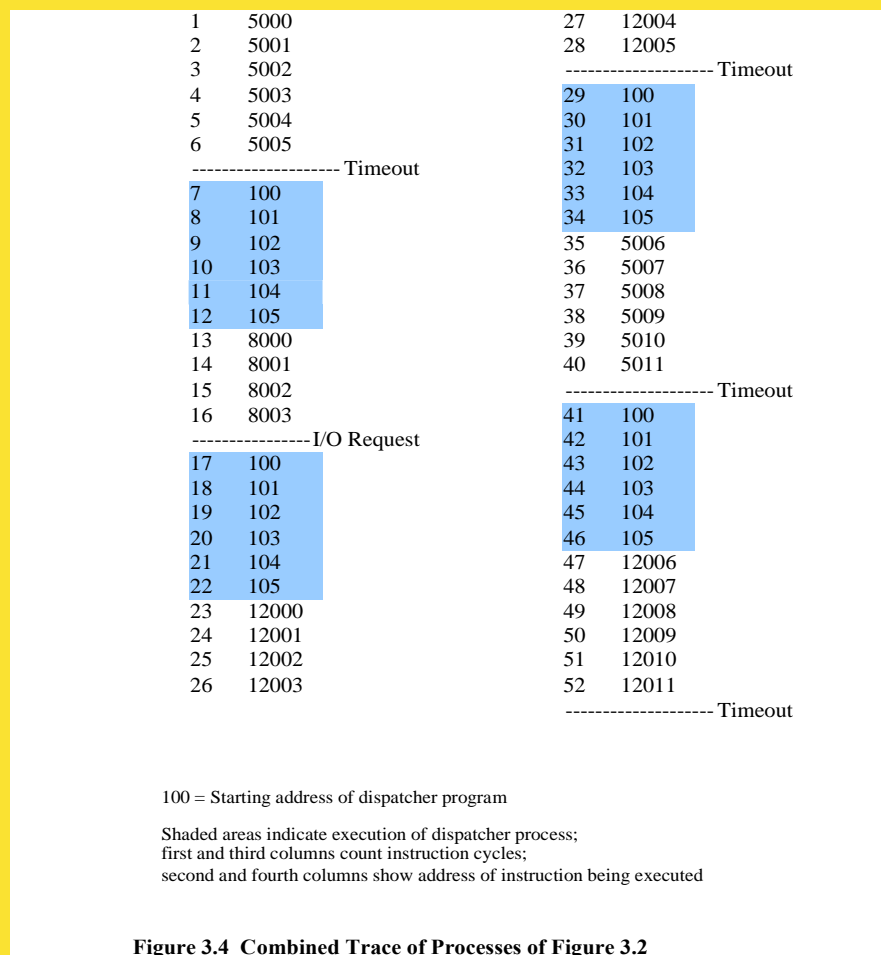


(b) Multiple blocked queues

Processes may have made I/O requests of different priorities (see last lecture, where interrupts can have different priorities) – each queue is for a given priority level

Figure 3.8 Queuing Model for Figure 3.6

# Recall Interleaving Processes A, B and C



**Figure 3.4 Combined Trace of Processes of Figure 3.2**

# Five-State Process Model: Interleaving Processes


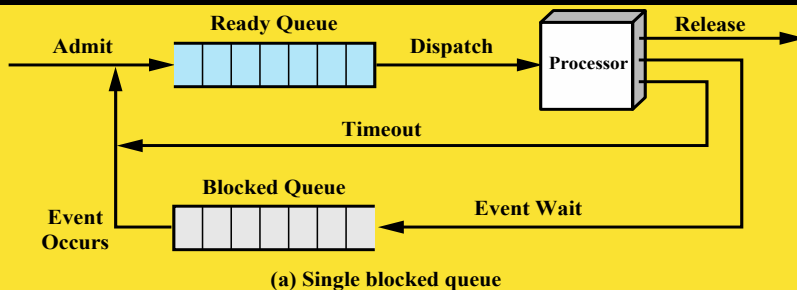


Figure 3.7 Process States for Trace of Figure 3.4

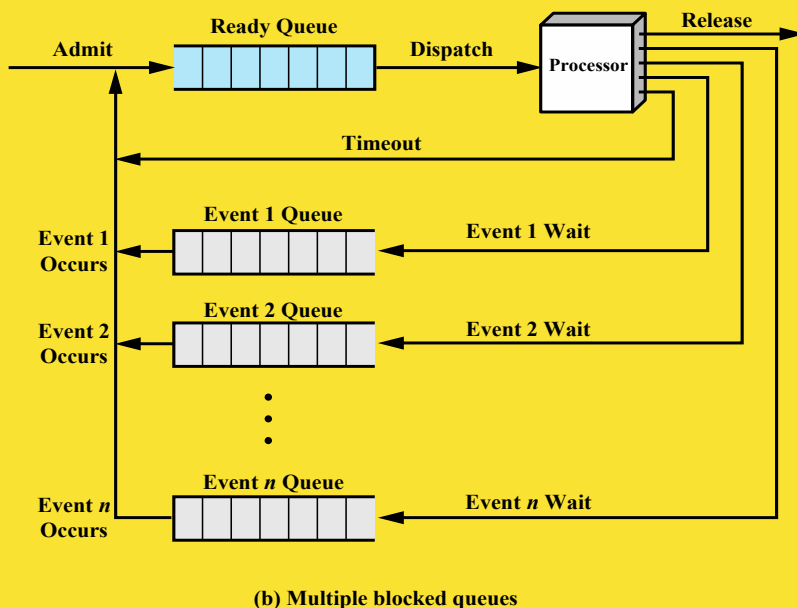
# Question

- How is the limited amount of main memory shared among many interleaved processes?

# But ...



CPU's are very fast, so these queues could grow very quickly! With limited main memory, this could be a challenge.



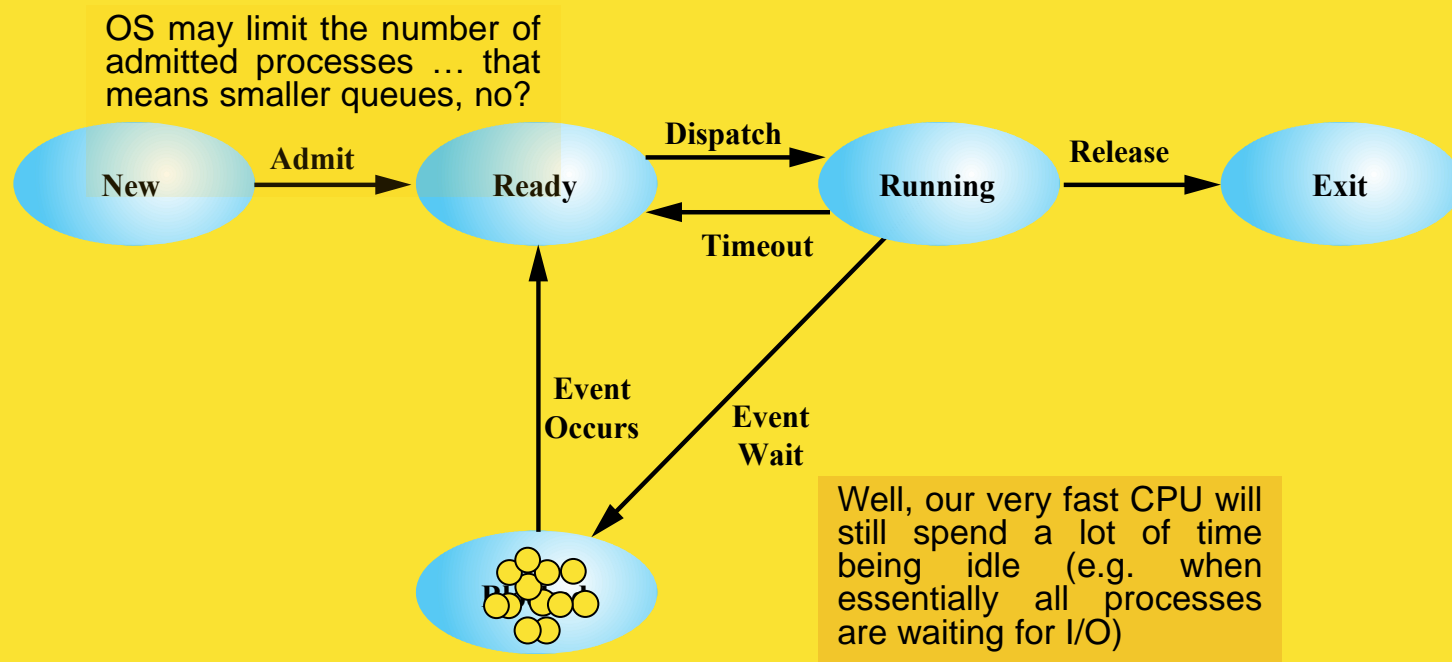
Why isn't having more main memory a good solution for this?

Ans:

- extra main memory comes with extra cost making the system less affordable
- In any case, new programs tend to require more and more main memory, so eventually the same problem of full queues will arise

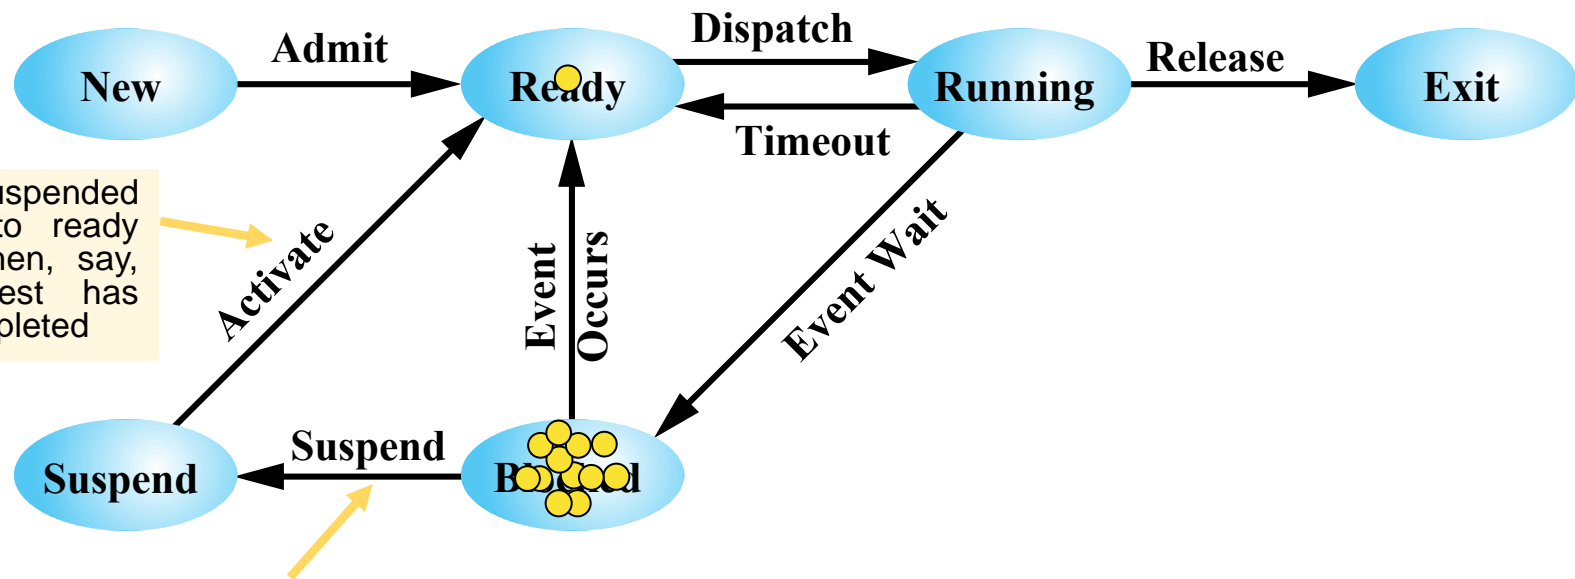
Figure 3.8 Queuing Model for Figure 3.6

# And, But ...



**Figure 3.6 Five-State Process Model**

# Six-State Process Model



Move suspended process to ready queue when, say, I/O request has been completed


Move part or all of some blocked processes to the hard disk

(a) With One Suspend State



# But ...

What happens if the ready queue becomes full (e.g. if, by servicing many new high priority users all at once reaches, the OS reaches its limits on the number of ready processes)?



Could there be other reasons for a process to be taken out of main memory temporarily? E.g. directly from the running state?

(a) With One Suspend State

Figure 3.9 Process State Transition Diagram with Suspend States

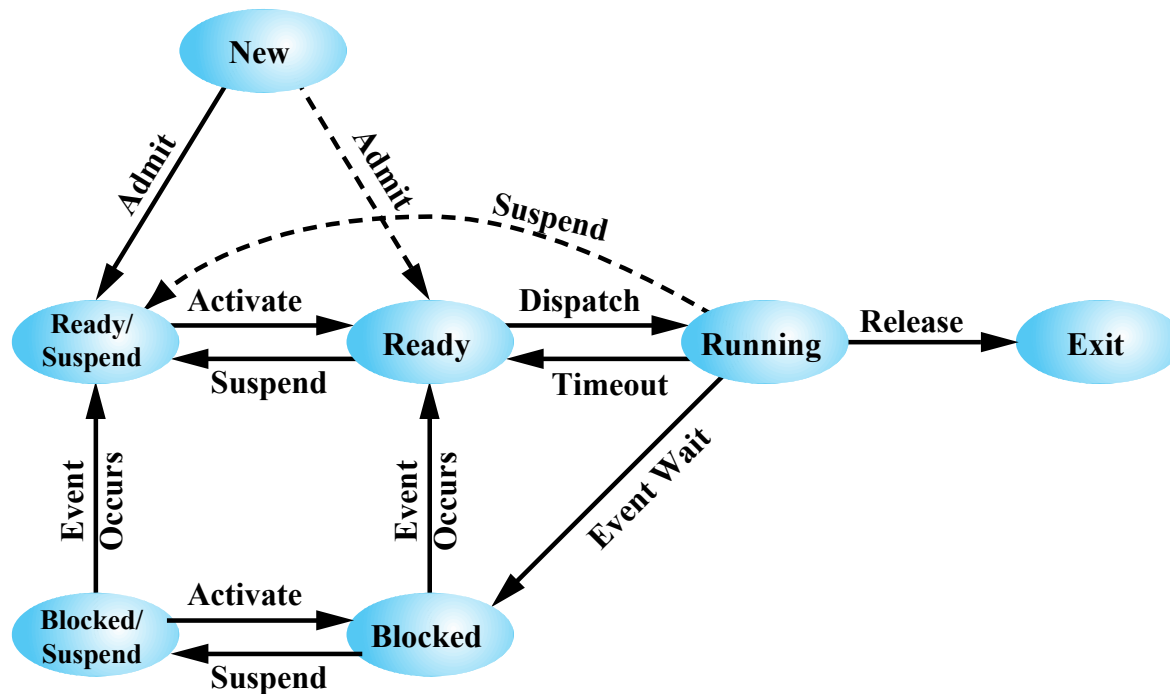
# Reasons for Entering Suspend State

Swapping	The OS needs to release sufficient main memory to bring in a process that is ready to execute.
Other OS reason	The OS may suspend a background or utility process or a process that is suspected of causing a problem.
Interactive user request	A user may wish to suspend execution of a program for purposes of debugging or in connection with the use of a resource.
Timing	A process may be executed periodically (e.g., an accounting or system monitoring process) and may be suspended while waiting for the next time interval.
Parent process request	A parent process may wish to suspend execution of a descendent to examine or modify the suspended process, or to coordinate the activity of various descendants.

# Qualities of a Process in the Suspend State

- A suspended process
  - is not immediately available for execution
  - is placed in a suspended state by a suspending agent
    - Itself (e.g. debugging), a parent process (e.g. for synch), or the OS
  - may not be waiting for an event (e.g. an I/O completion)
  - May not be removed from this state except by the suspending agent

# Seven-State process model



(b) With Two Suspend States

Figure 3.9 Process State Transition Diagram with Suspend States

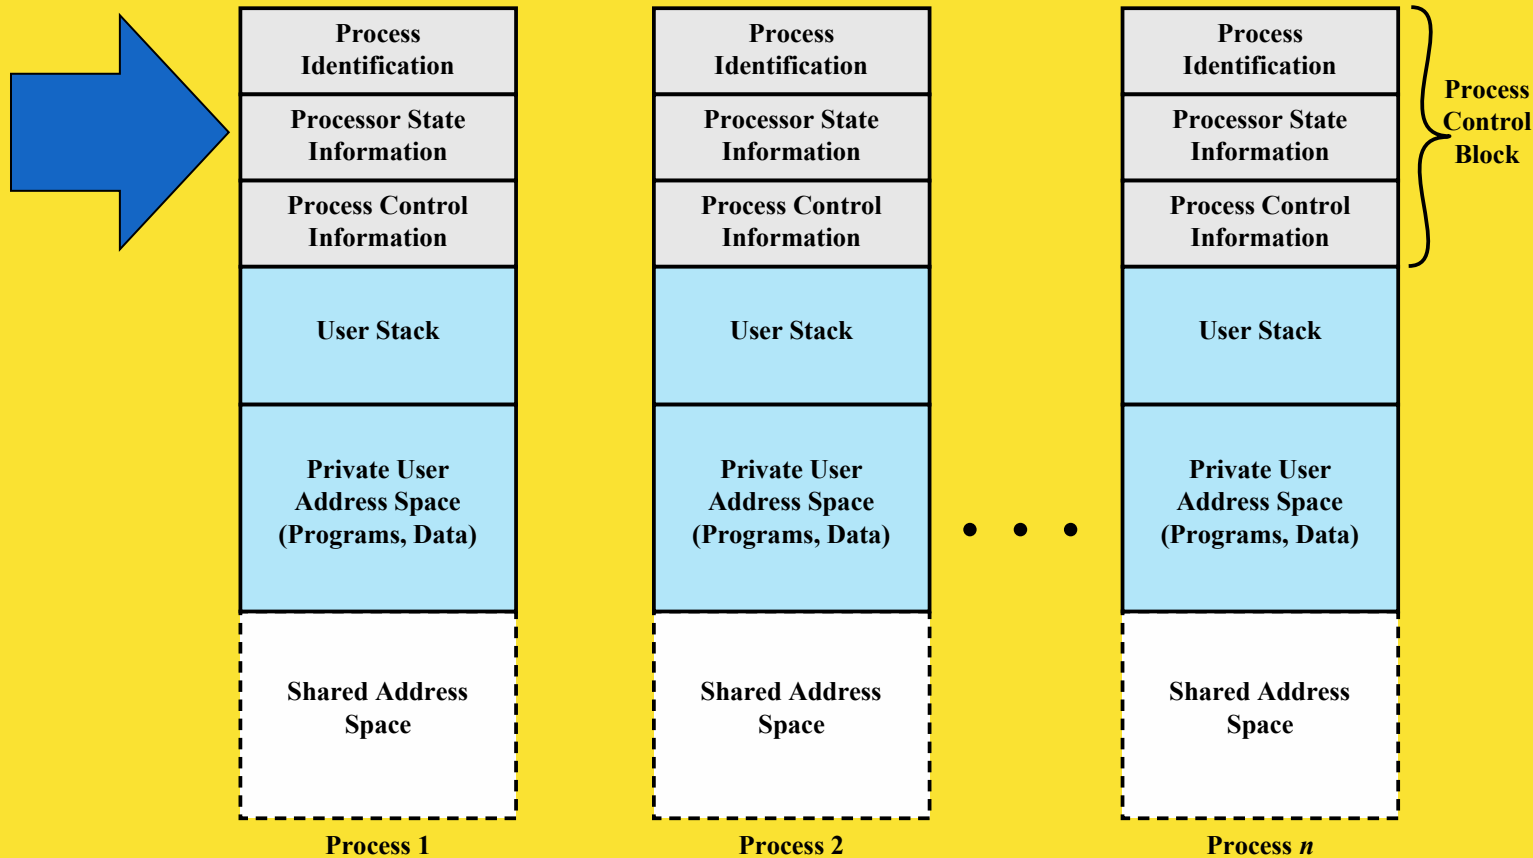
# Questions

- What does a process look like to the OS?
- How does an OS keep track of a process – the process's state, attributes, and the memory locations assigned to the process?
  - ans : special data structures in RAM and on the hard disk are used by the OS to do this
- What are these data structures?

# What does a process look like to the OS?

- A process is allocated a portion of main memory called the **process image**. This includes:
  - sufficient memory to hold the code and program data for the process
  - a **call stack**, used to keep track of a process's procedure calls, and parameter passing between procedures
  - memory for a special data structure called the **process control block** (PCB), that contains the process's attributes. These attributes are used by the OS for process control
- the precise location of the process image (in main memory) will depend on the memory management scheme used by the OS

# Process Images in Memory



# The Role of the Process Control Block (PCB)

- The most important data structure in an OS
  - Contains all of a process's information needed by the OS
  - PCBs are read and modified by virtually every module in the OS
  - The collection of all PCBs define the OS state
- Access to PCBs is easy: use process ID
- Protecting PCBs from being damaged is more challenging
  - A bug in a single OS routine could damage a PCB
  - A design change in the structure or semantics of the PCB could affect a number of modules within the OS

Process  
Identification

Processor State  
Information

Process Control  
Information



# Process Control Block (PCB)

- The OS maintains information about a process's attributes in PCB
- It resides in protected parts of memory
- Consists of 3 sections:
  - Process identification
  - Processor state information
  - Process control information

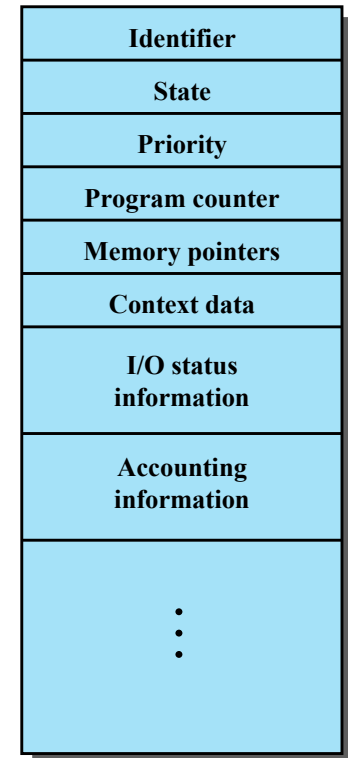

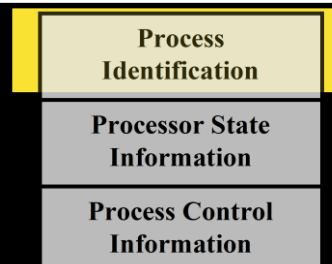


Figure 3.1 Simplified Process Control Block



# PCB Section : Process Identification

- The **process ID** is a unique number identifying the process
  - an index into the **primary process table**
  - part of a mapping the OS uses to locate process information in control tables
    - **Memory tables**
    - **I/O and file tables**
- When processes communicate, the process ID informs the OS of the destination of a particular communication
- When a process creates another process, process IDs distinguish the parent process from its child processes



# PCB Section : Processor State Information

The  
contents  
of  
processor  
registers

- User-visible registers
- Control and status registers
- Stack pointers

**Program  
status  
word  
(PSW)**

- Contains condition codes plus other status information
- EFLAGS register is an example of a PSW used by any OS running on an x86 processor

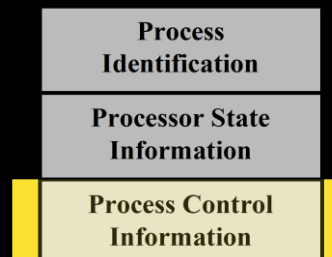
Process  
Identification

Processor State  
Information

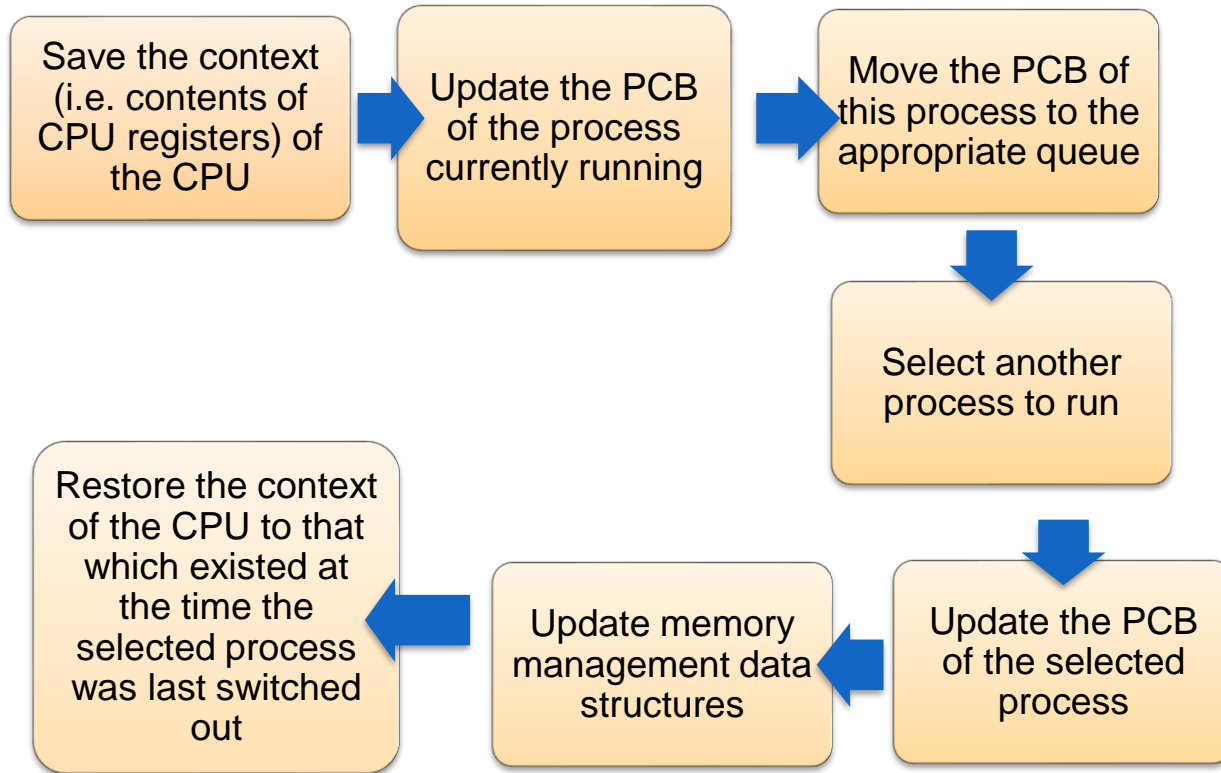
Process Control  
Information

# PCB Section : Process Control Information

- Used by the OS to assign the CPU to processes
  - Process state
  - Priority
  - Scheduling-related information
  - Identity of any event the process may be awaiting
- Links/pointers to other processes
- Interprocess communication
- Process privileges
- Memory management
- Resources assigned to the process



# OS re-assigning the CPU : Context Switch using PCBs



What happens when an OS re-assigns the CPU to a different process?

# OS Control Structures

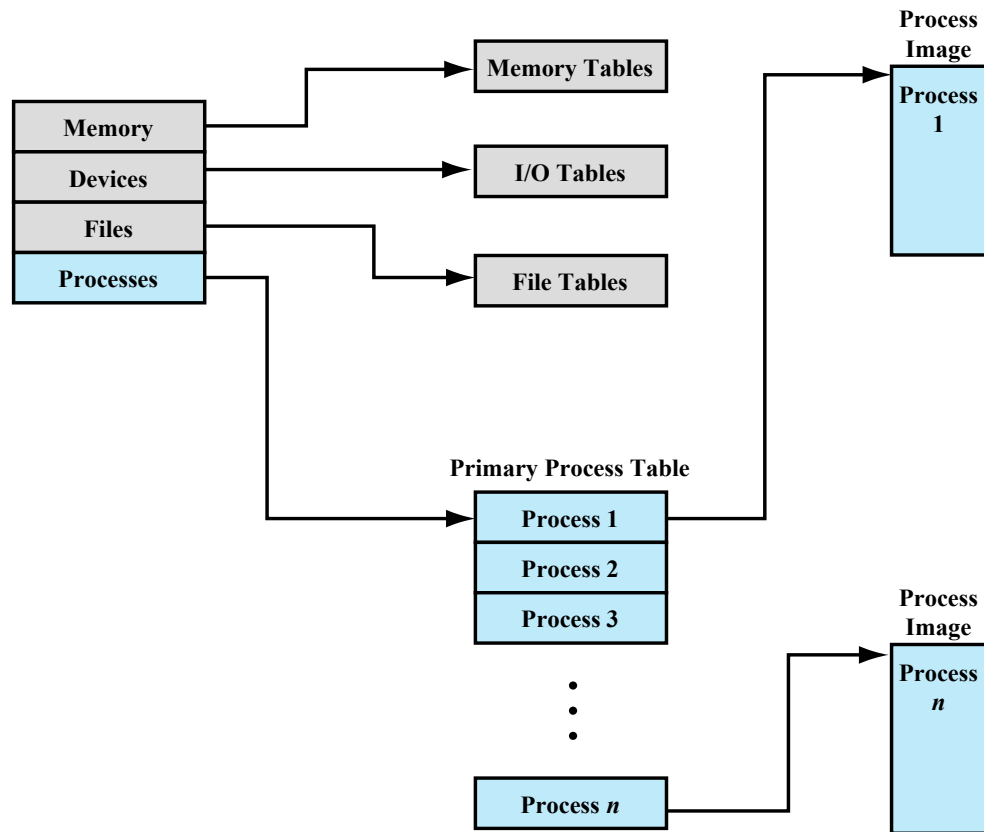


Figure 3.11 General Structure of Operating System Control Tables

# Memory Tables

- Used to keep track of both main and secondary memory
- Processes are maintained on secondary memory (i.e. hard-disk) using **virtual memory** or some swapping mechanism

## Must include information on:

Allocation of main memory to processes

Allocation of secondary memory to processes

Protection attributes of blocks of main or virtual memory

Information needed to manage virtual memory



# I/O Tables

- Used to manage I/O devices and channels
- At any given time, an I/O device may be either available or assigned to a particular process

If an I/O operation is in progress, the OS needs to know:

- The status of the operation
- The main memory location being used as the source or destination of the I/O transfer





# File Tables


These tables provide information about:

- Existence of files
- Location on secondary memory
- Current status
- Other attributes



# Process Tables

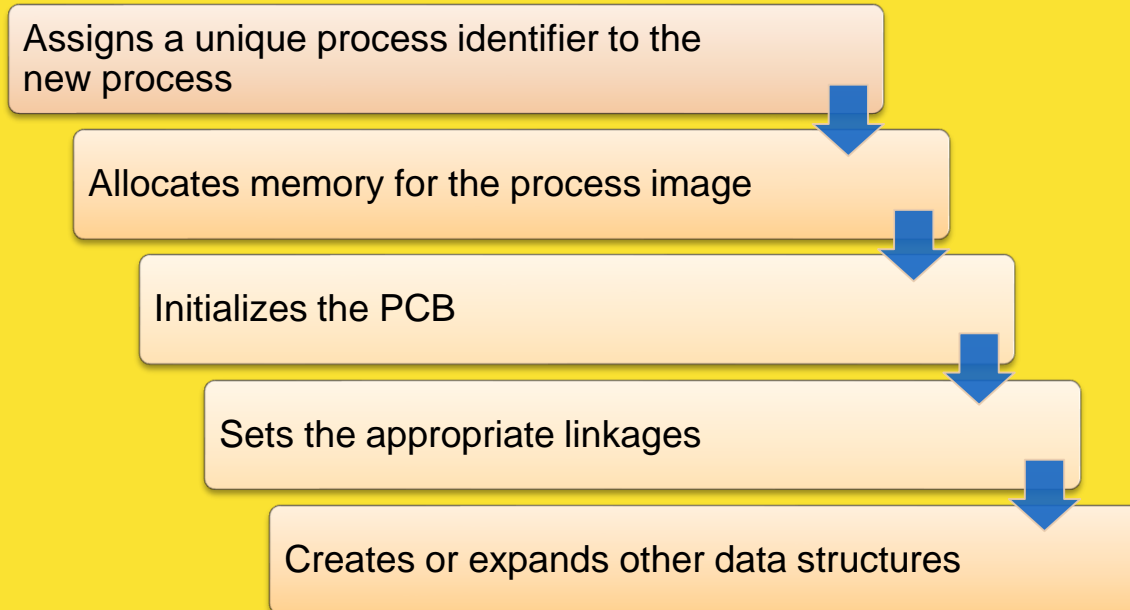
- Maintained by OS to manage processes
- Must include references to memory, I/O, and file tables, either directly or indirectly
- These tables must be accessible by the OS and, therefore, are subject to memory management



# Questions

- What does an OS do when it creates a process?
- Is the OS a process?


# Creating Processes



What does an OS do when it creates a process?


# Is the OS a Process?

No, it is completely **separate** from all processes




(a) Separate kernel

No, OS routines run **within** the address space of each running process



(b) OS functions execute within user processes



(c) OS functions execute as separate processes

No, it is **several** processes, running concurrently and performing different tasks

Figure 3.15 Relationship Between Operating System and User Processes